



Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:

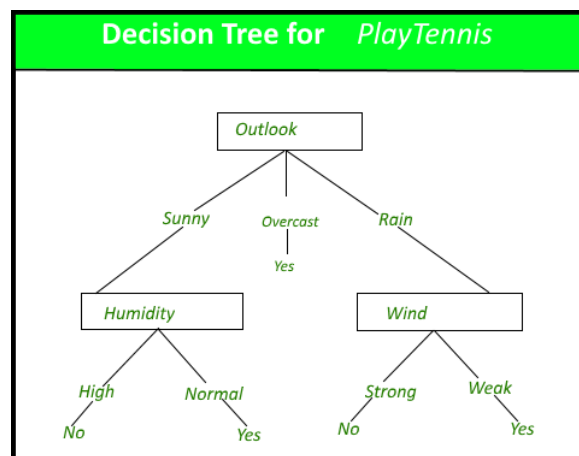


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:



>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op- Inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.



native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

In conclusion the Decision Tree model showed promising results on the Adult Census Income Dataset. Categorical attributes were properly handled through one-hot encoding, and data preprocessing which include replacing missing value with NaN, dropping tables, separating columns etc.

Hyperparameter tuning is a crucial step in optimizing the performance of a Decision Tree model because it allows to control the model's complexity by setting limits on some parameters. But to further improve the model's performance, by tuning hyperparameters like max depth, min samples split, etc., using techniques like Grid Search or Random Search can be achieved.

- Accuracy: Achieved an accuracy of 0.830, indicating that around 83% of predictions were correct.
- Precision: Precision of 0.652 suggests that among the instances predicted as positive, about 65.2% were actually positive.
- Recall: Recall of 0.694 indicates that the model captured around 69.4% of actual positive instances.
- F1 Score: The F1 score of 0.672 is the harmonic mean of precision and recall, offering a balanced assessment of the model's performance.

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import OneHotEncoder

data=pd.read_csv('adult.csv')
print(data)
```



	age	workclass	fnlwgt	education	education.num	marital.status	\
0	90	?	77053	HS-grad	9	Widowed	
1	82	Private	132870	HS-grad	9	Widowed	
2	66	?	186061	Some-college	10	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
...
24913	61	Private	477209	7th-8th	4	Married-civ-spouse	
24914	32	Private	70985	Assoc-voc	11	Married-civ-spouse	
24915	35	Private	241998	Bachelors	13	Married-civ-spouse	
24916	28	Private	249541	Some-college	10	Married-civ-spouse	
24917	57	Private	135339	Bachelors	13	Married-civ-spouse	

	occupation	relationship	race	sex	\
0	?	Not-in-family	White	Female	
1	Exec-managerial	Not-in-family	White	Female	
2	?	Unmarried	Black	Female	
3	Machine-op-inspct	Unmarried	White	Female	
4	Prof-specialty	Own-child	White	Female	
...
24913	Farming-fishing	Husband	White	Male	
24914	Machine-op-inspct	Husband	White	Male	
24915	Exec-managerial	Husband	White	Male	
24916	Craft-repair	Husband	White	Male	
24917	Exec-managerial	Husband	Asian-Pac-Islander	Male	

	capital.gain	capital.loss	hours.per.week	native.country	income
0	0	4356	40.0	United-States	<=50K
1	0	4356	18.0	United-States	<=50K
2	0	4356	40.0	United-States	<=50K
3	0	3900	40.0	United-States	<=50K
4	0	3900	40.0	United-States	<=50K
...
24913	0	0	54.0	United-States	<=50K
24914	0	0	40.0	United-States	<=50K
24915	0	0	45.0	United-States	>50K
24916	0	0	40.0	United-States	<=50K
24917	0	0	NaN	NaN	NaN

[24918 rows x 15 columns]

```
data.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week	
count	24918.000000	2.491800e+04	24918.000000	24918.000000	24918.000000	24917.000000	
mean	38.808813	1.897224e+05	10.130909	1408.191829	114.082190	40.546655	
std	13.676089	1.051236e+05	2.570755	8414.709871	457.306373	12.305727	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000	
25%	28.000000	1.180328e+05	9.000000	0.000000	0.000000	40.000000	
50%	37.000000	1.783190e+05	10.000000	0.000000	0.000000	40.000000	
75%	48.000000	2.373645e+05	13.000000	0.000000	0.000000	45.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000	

```
data.isnull().sum()
```

age	0
workclass	0
fnlwgt	0
education	0
education.num	0
marital.status	0
occupation	0
relationship	0
race	0
sex	0
capital.gain	0
capital.loss	0

```

hours.per.week    1
native.country    1
income            1
dtype: int64

# Replace '?' with NaN in the dataset
data.replace('?', pd.NA, inplace=True)

# Drop rows with missing values
data.dropna(inplace=True)

# Separate features and target
x = data.drop('income', axis=1)
y = data['income']

print(x)
print(y)

   age workclass  fnlwgt  education  education.num  marital.status \
1    82   Private  132870    HS-grad             9         Widowed
3    54   Private  140359    7th-8th             4         Divorced
4    41   Private  264663  Some-college          10         Separated
5    34   Private  216864    HS-grad             9         Divorced
6    38   Private  150601     10th             6         Separated
...   ...   ...   ...   ...   ...   ...
24912  44   Private  159580     12th             8         Divorced
24913  61   Private  477209    7th-8th             4  Married-civ-spouse
24914  32   Private   70985  Assoc-voc          11  Married-civ-spouse
24915  35   Private  241998  Bachelors          13  Married-civ-spouse
24916  28   Private  249541  Some-college          10  Married-civ-spouse

   occupation  relationship  race  sex  capital.gain \
1  Exec-managerial  Not-in-family  White  Female      0
3  Machine-op-inspct  Unmarried  White  Female      0
4  Prof-specialty  Own-child  White  Female      0
5  Other-service  Unmarried  White  Female      0
6  Adm-clerical  Unmarried  White  Male      0
...   ...   ...   ...   ...   ...
24912  Transport-moving  Not-in-family  White  Female      0
24913  Farming-fishing  Husband  White  Male      0
24914  Machine-op-inspct  Husband  White  Male      0
24915  Exec-managerial  Husband  White  Male      0
24916  Craft-repair  Husband  White  Male      0

   capital.loss  hours.per.week  native.country
1           4356           18.0  United-States
3           3900           40.0  United-States
4           3900           40.0  United-States
5           3770           45.0  United-States
6           3770           40.0  United-States
...   ...   ...   ...
24912           0           40.0  United-States
24913           0           54.0  United-States
24914           0           40.0  United-States
24915           0           45.0  United-States
24916           0           40.0  United-States

[23108 rows x 14 columns]
1      <=50K
3      <=50K
4      <=50K
5      <=50K
6      <=50K
...
24912  <=50K
24913  <=50K
24914  <=50K
24915  >50K
24916  <=50K
Name: income, Length: 23108, dtype: object

# Separate categorical and numerical columns
categorical_columns = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
numerical_columns = ['age', 'fnlwgt', 'education.num', 'capital.gain', 'capital.loss', 'hours.per.week']

x_categorical = x[categorical_columns]
x_numerical = x[numerical_columns]

# Apply one-hot encoding to categorical features
encoder = OneHotEncoder()
x_categorical_encoded = encoder.fit_transform(x_categorical)

```

```
# Split the data into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(x_encoded, y, test_size=0.3, random_state=1)

# Apply SMOTE to balance the class distribution
smote = SMOTE()
x_train_resampled, y_train_resampled = smote.fit_resample(x_train, y_train)

# Create Decision Tree classifier object
clf = DecisionTreeClassifier(criterion='entropy', min_samples_split=8, max_depth=10)


# Train Decision Tree Classifier
clf.fit(x_train_resampled, y_train_resampled)

# Predict on the test set
predictions = clf.predict(x_test)

accuracy = accuracy_score(y_test, predictions)
print("Accuracy:", accuracy)

Accuracy: 0.83888648492716
```

```
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier, plot_tree
plot_tree(clf)
plt.show()
```



```
# Evaluate the model
accuracy = accuracy_score(y_test, predictions)
precision = precision_score(y_test, predictions, pos_label='>50K')
recall = recall_score(y_test, predictions, pos_label='>50K')
f1 = f1_score(y_test, predictions, pos_label='>50K')

# Print evaluation metrics
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)

Accuracy: 0.83888648492716
Precision: 0.696604600219058
Recall: 0.6931880108991826
F1 Score: 0.6948921059819722
```

✓ 11s completed at 11:46 PM

● ×